

What we're doing

Aggregate finished contacts-v1 validation-loss runs from #75, #117, and #146 into one normalized, auditable scaling dataset.

Why

The sweeps span model sizes and training budgets but use different metric keys, tag schemas, and sweep subversion conventions. A clean table is the prerequisite for comparing parameter and token scaling without mixing obsolete runs.

Results so far

The first fetch contains 129 finished latest-subversion runs: 63 from #75, 50 from #117, and 16 from #146. Every row has the requested parameters, tokens, epochs, weight decay, learning rate, batch size, and validation loss. The fetcher records the exact loss key and provenance of each normalized field.

The first view plots every run against epoch count. Rings mark the lowest loss within each model-size/epoch group; issue-specific marker shapes keep the three sweeps distinguishable. Losses above 3.2 remain represented at the upper edge with small upward carets.

validation_loss_vs_epochs.png

All 129 finished runs; color = model size, marker = issue, rings = best per size/epoch, carets = loss above 3.2.



```
$ plot_validation_loss_by_epochs.py
```