

What we're doing

Optimize the 1.5B no-crops contacts-v1 model trained on AFDB and ESM-Atlas data. The first milestone is a ten-step DCLM nano smoke test of MarInFold's package-only MarIn training path on a ``v6e-4`` in ``europe-west4``.

Why

The protein sweep is expensive. A minimal run first verifies Iris dispatch, the packaged Marin stack, regional storage routing, TPU availability, and checkpoint logging without coupling the experiment to the sibling Marin source checkout.

Data safety

The smoke script adopts the completed regional Llama-3-tokenized DCLM cache as an immutable external artifact. It contains no tokenization recipe, so missing data causes failure rather than retokenization.

Results so far

The DCLM preflight passed (1,024 complete Llama-3-tokenized shards), and the graph contained only an adopted cache plus training. Iris launched ``run_levanter_train_lm-ad08751d``; the runtime confirmed a ``v6e-4`` and ``auto_build_caches=False``. The run stopped before step 0 because the authenticated account lacks Models write access to ``open-athena``. Grant that permission, then rerun unchanged.