

What we're doing

RL on #230's ``<contacts-v1.multi>`` checkpoint, with the reward computed over the **sections of a single rollout** rather than over the rollouts of a group.

Four arms, all warm-started from ``plm-exp230-cv1-multi-1_5b-.../step-1988``: **M-C** rewards each section by its leave-one-out contribution to its own rollout's consensus; **M-F** rewards the last section's F1; **M-B** rewards the best section's F1 (an oracle ceiling); **M-0** is a zero-LR control.

Why

#208 ran eleven scored runs across five reward designs and none improved consensus R-precision. The mechanism was a **unit mismatch**: the reward acted on one rollout, while the metric scored a vote over 100 **independent** rollouts — an object no single rollout can see. Making each rollout individually better made the hundred redundant, and consensus scoring cannot rank a pair that no rollout emits.

Under `<contacts-v1.multi>` the candidate set lives **inside one rollout** — ~22 contact sets in one sequence. A reward on the aggregate of one rollout's sections is therefore computed on the same kind of object the metric scores, and its credit assignment is within the sequence, where the gradient can reach it.

That is the only reason to expect a different outcome from #208, and it is the one thing this experiment tests.

The risk, preregistered

#230's checkpoint already sits at mean pairwise Jaccard **0.304** between its sections, past exp200's 0.30 diversity-collapse criterion **before any RL**. #208's dominant failure was RL collapsing diversity. So union coverage, total votes and votes-per-pair are reported every batch, and three diversity gates are kill criteria rather than diagnostics.

Results — the hypothesis is half right

Moving the reward's unit from *rollouts of a group* to *sections of one rollout* produces what #208 could not: an RL checkpoint that ****improves consensus R-precision****. Arm M-C step-18 is the best on every aggregation mode — consensus 0.5750 (+0.0077), oracle-best 0.5578 (+0.0235), last 0.5267 (+0.0701). On eval2-natural, last goes 0.1696 → 0.2421.

It does not close the gap to independent sampling. 22 plain rollouts read 0.5896 against 0.5750 — the closest a multi-mode number has come, and 0.0146 short.

#208's result is a dose-response, not a verdict

Consensus against distance moved: 0.5673 at KL 0, **0.5750 at 0.007**, 0.5741 at 0.016, 0.5529 at 0.031, 0.3969 at 0.486. Every reward here helps at small KL and damages at large. #208 ran its arms at KL 0.06–0.10 and to 3.96 — past the peak on all of them — and its two arms under 0.0015 never moved at all. The window it needed lay **between the two learning rates it tried**.

Reward shape decides how a run fails, not whether

M-C and M-F halve the contacts emitted while becoming *more* diverse (Jaccard 0.60× and 0.44×). M-B holds the volume and emits the same contacts 1.7× as often. Both routes end at the same coverage floor. #208 found these two modes across different reward families; here they are produced deliberately by reward shape, on one model and one data order.


Two lessons for the next reward

**** $E[r] = p - \bar{p}$ is necessary and not sufficient.**** M-C's advantage is centred so $E[A] = 0$ holds exactly — and it still shrank the policy, because 45 % of section marginals are an atom at exactly zero, so the *median* section is negative while the mean is zero. Checkable from a histogram before the run.

****Gate on union/R , and on precision.**** The preregistered coverage gate stopped all three arms; union/R never left 2.8–4.6 in any of them, including the collapsed one, whose union/R was *higher* than the warm start's. Coverage was never binding. Precision (0.50 → 0.14) was.


coverage_vs_kl.png

Union coverage against distance moved. #208 fitted, and then refuted, a model in which diversity loss depends only on how far the policy travels. It does not hold here either: M-C crosses the 80% floor by KL 0.013, M-B around 0.016, and M-F not until 0.036 -- so per unit of KL the three rewards cost very different amounts of coverage. They differ in HOW they lose it (volume for M-C and M-F, redundancy for M-B) and in how fast, but not in WHETHER.



gates_over_training.png

The three preregistered diversity gates, per training batch, rolling median over 6 batches. Jaccard is the panel that separates the arms: M-B (best-section reward) rises past #230's 0.304 -- it pays to emit your best guess repeatedly -- while M-C and M-F fall, i.e. their sections became MORE complementary. All three nonetheless hit the union coverage floor and were stopped. Faint lines are raw batches; that swing is the protein draw, and the zero-LR control reproduced it with a policy that did not change.



```
$ make_plots.py --steps data/training_steps.csv.gz --out plots/
```