

Question and design

Can sequence-only inference return both observed folds of a fold-switching protein? We ran the exp277 contacts-v1 checkpoint on 67 non-capped pairs. The primary test set has 29 pairs with at least 98% identity, held apart by sequence group from 15 development pairs. Seventeen of those 29 are literally identical. Every method returned at most 16 ranked maps chosen without reference contacts or structures.

The primary comparison is negative

At 100 shared root plus 100 arm rollouts per protein, 10-contact clustered branching improved mean minority-fold enrichment by only 0.013 over independent sampling. The paired 95% interval is -0.014 to +0.043; mean minority recall was essentially unchanged. Blind dual-contact hits were 3/29 versus 2/29. The extra hit failed a 3D check. On the 17 exact-sequence test pairs, both methods had one contact hit and the paired interval still spans zero. We cannot claim a reliable search improvement.

What was actually discovered


For 3j7w/3j7v, both independent and branching top-16 shortlists included maps that reconstructed distinct switching-region structures. With the Fold1 sequence input, region-touch IDDT against (Fold1, Fold2) was (0.477, 0.206) and (0.356, 0.702) for the branching shortlist; the independent shortlist also favored the respective folds. True-contact controls showed comparable separation. This is an existence proof for blind inference discovery, not a branching advantage.

Where the search is limited

Four of seven branch-10 oracle dual-contact hits were present among 200 candidates but missed by the blind 16-map selector. In a separate plain-sampling run, blind dual hits were 2, 3, and 3 at 100, 200, and 500 rollouts; the oracle pool reached 9 by 200. The first priority is better reference-free ranking of coherent minority maps. Contact-level cutoffs are operational and not calibrated as structural thresholds; only four proteins received 3D checks.


budget_curve.png

A separate 500-rollout run, evaluated at its first 100, 200, and 500 candidates. The blind selector returns 16 maps at each budget; the oracle line is diagnostic and cannot be used at inference time.



dual_contact_hits.png


Bars count pairs meeting the prespecified contact-level dual-hit criterion among 16 blindly selected maps. Diamonds show the oracle ceiling in the same candidate pool, which cannot serve as an inference method.



helico_3j7_structural_check.png

Helico structures from two reference input chains; black points are the two inputs. Both 200-rollout shortlist methods yield fold-specific switching-region structures on this protein. True-contact controls show the attainable separation for this reconstruction model.

3j7w/3j7v: both blind selectors recover two local folds



paired_enrichment.png

Each point is one held-out primary fold-switching pair. Both methods use the same 100 root rollouts plus 100 arm rollouts; the diagonal is parity.

