

Sampling contacts from sequence

A figure-first writeup on existing predictors, oracle contacts, MarinFold and useful sampling diversity.

Fixed model: exp277 step 266344, 1.47B parameters, 248.584B raw training tokens. Natural proteins lead: 97 validation + 217 publication test. Designs remain separate. AF2/AF3/Boltz-2 comparators use shared archived MSAs and no templates, with confidence selection.

Predictor generation → cached analysis tables → fast static and interactive rendering. Every numerical figure element maps back to a source CSV row.

What the contact and confidence analyses show

AF2 / AF3 / Boltz-2 GDT-TS: 0.854 / 0.860 / 0.861 on 305 natural proteins. At MSA depth <10: 0.196 / 0.332 / 0.239, with only five proteins.

MarinFold contact R-precision: 0.561 overall; 0.345 at MSA depth <10 versus 0.617 at ≥ 1000 . Only five natural proteins occupy the shallowest tier.

Oracle pTM ranks 1, 5, 1, 1, 1 out of 101 maps for the five low-depth proteins. Highest pTM selects each map's Helico sample; no ipTM or clash penalty. 100 ESMFold2 maps/protein, a shared eligible-pair mask, three Helico samples/map. Original ESMFold2 median TM-score: 0.848, 0.860, 0.881, 0.649, 0.839. pTM-source TM-score Spearman: 0.127, 0.083, 0.254, 0.288, 0.613. Protein order: 8ii8_A, 8oxk_A, 8qoh_A, 8ux2_A, 8wrx_A. Only five biological examples.

Top-L contacts raise Helico GDT-TS from 0.150 to 0.504 on 305 matched natural proteins. The paired gain is 0.354 [0.323, 0.387]. All cuts were fixed before test inference.

Can more AlphaFold3 sampling recover accurate folds?

5,000 predictions: 1,000 independent full runs for each of five natural test proteins. One diffusion sample per seed, ten recycles, fixed archived MSA, no templates.

8oxk_A: best TM 0.516 in the original 25 → 0.818 in 100 fresh runs → 0.947 in 1,000. pTM selects the 0.947 structure (seed 10470, draw 471). 20 / 1,000 predictions reach TM ≥ 0.8 ; the first is draw 85. At budget 100, pTM selected a different structure with TM 0.643.

8qoh_A / 8ux2_A: no TM ≥ 0.8 in 1,000; best TM 0.502 / 0.635. 8ii8_A / 8wrx_A: already accurate; best TM 0.881 / 0.936.

Per-protein plots separate oracle best-of-N from pTM-selected TM. The first 100 runs are unchanged; original baseline scores reproduce. Five biological examples, fixed generation settings, no model tuning.

AF3 sampling in predictor and MSA-depth context

8oxk_A: ESMFold2 already reaches TM 0.940; AF3 at 1,000 reaches 0.947. Five low-depth proteins, mean TM: ESMFold2 0.854; AF3 baseline 0.603; AF3 at 1,000: 0.681 with official ranking, 0.719 with pTM, 0.780 oracle best. TM ≥ 0.8 : ESMFold2 4/5; original AF3 2/5; expanded AF3 3/5.

At depths 10–99 / 100–999 / ≥ 1000 , baseline AF3 means: 0.891 / 0.932 / 0.955. ESMFold2: 0.691 / 0.874 / 0.944. Same 305 proteins; bin $n=5$ / 20 / 60 / 220. Different proteins occupy each bin; extended sampling only exists at depth < 10 .

Official defaults: ten recycles, five diffusion samples per user-specified seed. Our baseline: 5 full seeds x 5 samples; expanded: 1,000 full seeds x 1 sample. Both use fixed archived MSAs and no templates, not the native search pipeline. 40x structures, 200x trunk evaluations; neither is a measured runtime ratio. Official ranking includes disorder/clash terms; pTM ranking is shown separately.

Useful diversity is the open question

314 proteins $\times$ 100 samples: consensus R-precision 0.561; oracle best individual map 0.528. Paired difference -0.032 , 95% bootstrap interval $[-0.038, -0.025]$.

All 100 maps are distinct for every protein; mean Jaccard 0.277, true-contact union recall 0.943. This is not an absence of contact-set diversity, and it does not directly measure distinct folds.

Next: better whole-map candidates, inference-time search and post-training.

Low benchmark MSA depth versus ESMC training sources

ESMC documents UniRef 2023_02, MGnify 2023_02 and JGI (July 2023), not Logan. Low retrieved alignment depth does not establish novelty to the language model.

New historical UniRef search, with MSA filtering disabled: Significant hits covering at least half the query: 1, 4, 1, 15, 1. Protein order: 8ii8_A, 8oxk_A, 8qoh_A, 8ux2_A, 8wrx_A. 14 of the 15 well-covered 8ux2_A hits are 97.9–100% identical to the query. Counts exclude the added query; they are not Neff or final training-cluster counts.

Online MGnify30 hits count heavily clustered representatives, not training-source depth. ESMC uses 70%-identity clustering; historical MGnify90 is the closer public search target. Full historical MGnify and JGI are not yet searched. Complete training depth is unknown. See TRAINING_SOURCE_SEARCH.md and data/training_source_depths.csv for details. Offline reduction: `uv run python generation/search_training_sources.py prepare`

Precision at L/5 and the KNN comparison


$k = \max(1, \text{floor}(L/5))$; L is the input sequence length, not the resolved-residue count. Same 314 natural proteins, nine predictors, candidate pairs and archived predictions. All-range / long-range P@L/5: Marifold 0.830 / 0.761; sequence-KNN 0.657 / 0.595. At depth <10 (five proteins): Marifold 0.595; KNN 0.102; ESMFold2 0.727. KNN transfers from ten native decontaminated-corpus neighbors; no redesign index.

Same 100-rollout pools: mean individual 0.400; consensus 0.829; oracle best 0.609. Oracle reselected by P@L/5. Emission order ranks individual maps; short maps retain k. Invalid individual maps receive zero. Consensus follows the original vote rules. All 5,994 predictor and 1,884 sampling R-precision cells reproduce before rescoreing.

Every static plot has native vector PDF/SVG and a 7,200-pixel-wide PNG. At 300 dpi the raster supports 24-inch-wide panels; vector PDFs scale further. POSTER.md links the public collection. Preprocessing remains separate from rendering.


01_predictors.png

Natural FoldBench monomers; the same 305 proteins in every structural arm. Points are protein means; bars are 95% protein-bootstrap intervals. MSA depth counts sequences, including the query, in the alignment used by Protenix-v2 + MSA. AlphaFold2/3 and Boltz-2 use the same archived MSA queries and alignments underlying these depth bins, with templates disabled and protein-chain-only inputs. Confidence selects among five AF2 pTM models (three recycles), 25 AF3 samples (five seeds, ten recycles), or 25 Boltz-2 samples (one seed, ten recycles). These are controlled shared-MSA runs, not the models' full default search pipelines.




01_predictors_lddt.png

Natural FoldBench monomers; the same 305 proteins in every structural arm. Points are protein means; bars are 95% protein-bootstrap intervals. MSA depth counts sequences, including the query, in the alignment used by Protenix-v2 + MSA. AlphaFold2/3 and Boltz-2 use the same archived MSA queries and alignments underlying these depth bins, with templates disabled and protein-chain-only inputs. Confidence selects among five AF2 pTM models (three recycles), 25 AF3 samples (five seeds, ten recycles), or 25 Boltz-2 samples (one seed, ten recycles). These are controlled shared-MSA runs, not the models' full default search pipelines.



01b_af3_sampling_8ii8_A.png


8ii8_A: each point is one independent AF3 trunk + diffusion run, seeds starting at 10000, ten recycles, fixed archived MSA and no templates. Full-precision pTM selects without using truth. Oracle best TM uses truth only for this diagnostic. Dashed line: prespecified TM ≥ 0.8 . Curves use the ascending-seed prefix; original dotted baseline selected among five seeds x five samples by official ranking_score. Data: af3_sampling_samples.csv and af3_sampling_curves.csv; original baseline: af3_sampling_original.csv. ESMFold2 baseline: af3_context_rows.csv. Official ranking_score selection is shown separately from pTM. All are in data/; preprocessing: prepare_af3_sampling.py and prepare_af3_context.py.



\$ render_af3_sampling.py

01b_af3_sampling_8oxk_A.png


8oxk_A: each point is one independent AF3 trunk + diffusion run, seeds starting at 10000, ten recycles, fixed archived MSA and no templates. Full-precision pTM selects without using truth. Oracle best TM uses truth only for this diagnostic. Dashed line: prespecified TM ≥ 0.8 . Curves use the ascending-seed prefix; original dotted baseline selected among five seeds x five samples by official ranking_score. Data: af3_sampling_samples.csv and af3_sampling_curves.csv; original baseline: af3_sampling_original.csv. ESMFold2 baseline: af3_context_rows.csv. Official ranking_score selection is shown separately from pTM. All are in data/; preprocessing: prepare_af3_sampling.py and prepare_af3_context.py.



\$ render_af3_sampling.py

01b_af3_sampling_8qoh_A.png


8qoh_A: each point is one independent AF3 trunk + diffusion run, seeds starting at 10000, ten recycles, fixed archived MSA and no templates. Full-precision pTM selects without using truth. Oracle best TM uses truth only for this diagnostic. Dashed line: prespecified TM ≥ 0.8 . Curves use the ascending-seed prefix; original dotted baseline selected among five seeds x five samples by official ranking_score. Data: af3_sampling_samples.csv and af3_sampling_curves.csv; original baseline: af3_sampling_original.csv. ESMFold2 baseline: af3_context_rows.csv. Official ranking_score selection is shown separately from pTM. All are in data/; preprocessing: prepare_af3_sampling.py and prepare_af3_context.py.



\$ render_af3_sampling.py

01b_af3_sampling_8ux2_A.png


8ux2_A: each point is one independent AF3 trunk + diffusion run, seeds starting at 10000, ten recycles, fixed archived MSA and no templates. Full-precision pTM selects without using truth. Oracle best TM uses truth only for this diagnostic. Dashed line: prespecified TM ≥ 0.8 . Curves use the ascending-seed prefix; original dotted baseline selected among five seeds x five samples by official ranking_score. Data: af3_sampling_samples.csv and af3_sampling_curves.csv; original baseline: af3_sampling_original.csv. ESMFold2 baseline: af3_context_rows.csv. Official ranking_score selection is shown separately from pTM. All are in data/; preprocessing: prepare_af3_sampling.py and prepare_af3_context.py.



\$ render_af3_sampling.py

01b_af3_sampling_8wrx_A.png


8wrx_A: each point is one independent AF3 trunk + diffusion run, seeds starting at 10000, ten recycles, fixed archived MSA and no templates. Full-precision pTM selects without using truth. Oracle best TM uses truth only for this diagnostic. Dashed line: prespecified TM ≥ 0.8 . Curves use the ascending-seed prefix; original dotted baseline selected among five seeds x five samples by official ranking_score. Data: af3_sampling_samples.csv and af3_sampling_curves.csv; original baseline: af3_sampling_original.csv. ESMFold2 baseline: af3_context_rows.csv. Official ranking_score selection is shown separately from pTM. All are in data/; preprocessing: prepare_af3_sampling.py and prepare_af3_context.py.



\$ render_af3_sampling.py


01c_af3_context.png

TM-score for the same five natural test proteins with MSA depth <10. Top block: archived predictor entries, including the 248B-token MarinFold top-L contacts through Helico. Middle: 100 or 1,000 fresh AF3 full runs, one diffusion sample per seed, selected by official ranking_score or pTM. Bottom: ground-truth diagnostics, oracle best AF3 TM and Helico conditioned on the oracle map. ESMFold2 is the archived baseline, not a selected member or average of the 100-contact-map decoy pool. AF3 baseline uses five seeds times five diffusion samples, ten recycles, shared archived MSAs and no templates. Expanded runs retain the same inputs and recycles. Compute budgets differ. Hover identifies the exact source CSV cell.




01d_af3_depth_context.png

Mean TM-score on the same 305 natural proteins as the main structural figures: 5, 20, 60 and 220 proteins across the four MSA bins. All archived predictors cover exactly this population; each protein receives equal weight. Hover shows 95% protein-bootstrap intervals and success counts; the prepared CSV retains every contributing row. The 1,000-run AF3 study covers only the five <10 proteins; other cells are not run. Stars mark ground-truth diagnostics. Different bins contain different proteins, so this does not estimate the causal effect of adding MSA sequences. All AF3 results use shared benchmark MSAs and no templates, not the native search pipeline. Budgets and training data differ between predictors.




02_oracle.png

*Oracle = ground-truth contacts AND non-contacts supplied to Helico without an MSA. This is an information upper bound, not a competing predictor or a matched-count true-contact experiment. Same 305 natural proteins as Figure 1. AlphaFold2/3 and Boltz-2 use the same archived MSA queries and alignments underlying these depth bins, with templates disabled and protein-chain-only inputs. Confidence selects among five AF2 pTM models (three recycles), 25 AF3 samples (five seeds, ten recycles), or 25 Boltz-2 samples (one seed, ten recycles). These are controlled shared-MSA runs, not the models' full default search pipelines.




02_oracle_lddt.png

*Oracle = ground-truth contacts AND non-contacts supplied to Helico without an MSA. This is an information upper bound, not a competing predictor or a matched-count true-contact experiment. Same 305 natural proteins as Figure 1. AlphaFold2/3 and Boltz-2 use the same archived MSA queries and alignments underlying these depth bins, with templates disabled and protein-chain-only inputs. Confidence selects among five AF2 pTM models (three recycles), 25 AF3 samples (five seeds, ten recycles), or 25 Boltz-2 samples (one seed, ten recycles). These are controlled shared-MSA runs, not the models' full default search pipelines.




02b_confidence.png

All five natural FoldBench proteins with MSA depth <10 (test split). Each dot is one of 100 single-sequence ESMFold2 predictions, converted to a full contact/non-contact map and folded by Helico; diamonds use the ground-truth oracle map. Every map has the same eligible-pair mask and three Helico diffusion samples. Highest pTM selects a sample per map; that pTM also ranks the 101 maps. No ipTM term, clash penalty or map filtering is used. Contact counts may vary. Rank 1 is highest pTM; ties are shown as rank intervals. Vertical jitter only separates points. These 500 decoys represent five proteins.




02c_accuracy_confidence_protein_8ii8_A.png

Highest pTM selects one of three Helico samples for every map; pTM also ranks the 101 maps. Left: original ESMFold2 TM-score (oracle reference = 1). Right: selected Helico reconstruction TM-score. Diamonds identify oracle contacts; correlations use only the 100 ESMFold2 maps. TM-score uses matched protein CA atoms; pTM retains all Helico input tokens. No ipTM or clash penalty.




02c_accuracy_confidence_protein_8oxk_A.png

Highest pTM selects one of three Helico samples for every map; pTM also ranks the 101 maps. Left: original ESMFold2 TM-score (oracle reference = 1). Right: selected Helico reconstruction TM-score. Diamonds identify oracle contacts; correlations use only the 100 ESMFold2 maps. TM-score uses matched protein CA atoms; pTM retains all Helico input tokens. No ipTM or clash penalty.




02c_accuracy_confidence_protein_8qoh_A.png

Highest pTM selects one of three Helico samples for every map; pTM also ranks the 101 maps. Left: original ESMFold2 TM-score (oracle reference = 1). Right: selected Helico reconstruction TM-score. Diamonds identify oracle contacts; correlations use only the 100 ESMFold2 maps. TM-score uses matched protein CA atoms; pTM retains all Helico input tokens. No ipTM or clash penalty.




02c_accuracy_confidence_protein_8ux2_A.png

Highest pTM selects one of three Helico samples for every map; pTM also ranks the 101 maps. Left: original ESMFold2 TM-score (oracle reference = 1). Right: selected Helico reconstruction TM-score. Diamonds identify oracle contacts; correlations use only the 100 ESMFold2 maps. TM-score uses matched protein CA atoms; pTM retains all Helico input tokens. No ipTM or clash penalty.




02c_accuracy_confidence_protein_8wrx_A.png

Highest pTM selects one of three Helico samples for every map; pTM also ranks the 101 maps. Left: original ESMFold2 TM-score (oracle reference = 1). Right: selected Helico reconstruction TM-score. Diamonds identify oracle contacts; correlations use only the 100 ESMFold2 maps. TM-score uses matched protein CA atoms; pTM retains all Helico input tokens. No ipTM or clash penalty.




03_method.png

Schematic; arrow widths do not encode amounts. Exp277 step 266,344: a scratch-trained 1.47B-parameter Qwen3, one epoch over 232,090,905 contact documents, totaling 248,583,762,834 raw tokens. Native and ProteinMPNN-redesigned sequences use AFDB / ESM Atlas source structures from the decontaminated corpus. Raw tokens differ from padded training slots.




04_contacts.png

Exp277 step 266,344, the 248.584B-raw-token model. All-range R-precision on 314 natural monomers (97 validation + 217 test); 100 resampled rollouts and vote ranking. R is the true-contact count in the resolved candidate universe. KNN indexes the native decontaminated corpus, not the additional redesign sequences. Only five natural proteins have depth <10. The menu also exposes long-range R-precision and separate split/design views. AlphaFold2/3 and Boltz-2 use the same archived MSA queries and alignments underlying these depth bins, with templates disabled and protein-chain-only inputs. Confidence selects among five AF2 pTM models (three recycles), 25 AF3 samples (five seeds, ten recycles), or 25 Boltz-2 samples (one seed, ten recycles). These are controlled shared-MSA runs, not the models' full default search pipelines.




04_contacts_long.png

Exp277 step 266,344, the 248.584B-raw-token model. All-range R-precision on 314 natural monomers (97 validation + 217 test); 100 resampled rollouts and vote ranking. R is the true-contact count in the resolved candidate universe. KNN indexes the native decontaminated corpus, not the additional redesign sequences. Only five natural proteins have depth <10. The menu also exposes long-range R-precision and separate split/design views. AlphaFold2/3 and Boltz-2 use the same archived MSA queries and alignments underlying these depth bins, with templates disabled and protein-chain-only inputs. Confidence selects among five AF2 pTM models (three recycles), 25 AF3 samples (five seeds, ten recycles), or 25 Boltz-2 samples (one seed, ten recycles). These are controlled shared-MSA runs, not the models' full default search pipelines.




04_contacts_pl5.png

P@L/5 on 314 natural FoldBench proteins, using $k=\max(1, \text{floor}(L/5))$ and L equal to the frozen input sequence length, not the number of resolved residues. Same nine predictors and same archived predictions as the R-precision panel, including AF2, AF3, Boltz-2 and decontaminated sequence-KNN. Ground truth and eligible pairs are unchanged: pyconfind degree ≥ 0.001 , resolved residues, separation ≥ 6 for all-range or ≥ 24 for long-range. Structure contacts rank by degree, MarinFold by votes; stable candidate-order ties. Means and 95% protein-bootstrap intervals. The menu exposes both ranges and split/design views. Preprocessing reproduces every archived R-precision score before using its P@L/5 counterpart.




04_contacts_pl5_long.png

P@L/5 on 314 natural FoldBench proteins, using $k=\max(1, \text{floor}(L/5))$ and L equal to the frozen input sequence length, not the number of resolved residues. Same nine predictors and same archived predictions as the R-precision panel, including AF2, AF3, Boltz-2 and decontaminated sequence-KNN. Ground truth and eligible pairs are unchanged: pyconfind degree ≥ 0.001 , resolved residues, separation ≥ 6 for all-range or ≥ 24 for long-range. Structure contacts rank by degree, MarinFold by votes; stable candidate-order ties. Means and 95% protein-bootstrap intervals. The menu exposes both ranges and split/design views. Preprocessing reproduces every archived R-precision score before using its P@L/5 counterpart.




04b_knn.png

Same 314 natural proteins and fixed 248B-token MarInFold checkpoint. Sequence-KNN transfers contacts from the ten nearest sequences in the native decontaminated training corpus; it does not index the additional ProteinMPNN redesign sequences. Points are protein means with 95% protein-bootstrap intervals. The menu exposes all-range and long-range P@L/5 and R-precision. Paired MarInFold-minus-KNN differences and their intervals are in pl5_paired_deltas.csv. This is a contact-prediction baseline, not a structure predictor.




04b_knn_long.png

Same 314 natural proteins and fixed 248B-token MarInFold checkpoint. Sequence-KNN transfers contacts from the ten nearest sequences in the native decontaminated training corpus; it does not index the additional ProteinMPNN redesign sequences. Points are protein means with 95% protein-bootstrap intervals. The menu exposes all-range and long-range P@L/5 and R-precision. Paired MarInFold-minus-KNN differences and their intervals are in pl5_paired_deltas.csv. This is a contact-prediction baseline, not a structure predictor.




04b_knn_r_precision.png

Same 314 natural proteins and fixed 248B-token MarInFold checkpoint. Sequence-KNN transfers contacts from the ten nearest sequences in the native decontaminated training corpus; it does not index the additional ProteinMPNN redesign sequences. Points are protein means with 95% protein-bootstrap intervals. The menu exposes all-range and long-range P@L/5 and R-precision. Paired MarInFold-minus-KNN differences and their intervals are in pl5_paired_deltas.csv. This is a contact-prediction baseline, not a structure predictor.




04b_knn_r_precision_long.png

Same 314 natural proteins and fixed 248B-token MarInFold checkpoint. Sequence-KNN transfers contacts from the ten nearest sequences in the native decontaminated training corpus; it does not index the additional ProteinMPNN redesign sequences. Points are protein means with 95% protein-bootstrap intervals. The menu exposes all-range and long-range P@L/5 and R-precision. Paired MarInFold-minus-KNN differences and their intervals are in pl5_paired_deltas.csv. This is a contact-prediction baseline, not a structure predictor.




05_folding.png

The 248B-token model through Helico: request the top-L predicted contacts, generate three diffusion samples, select by highest Helico confidence at that fixed cut; no MSA. Helico drops pairs closer than six residues after coordinate mapping; requested and effective counts are saved. Same 305 natural proteins in every arm (95 validation + 210 test). GDT-TS and IDDT are separate metrics. See coverage.csv for exclusions. AlphaFold2/3 and Boltz-2 use the same archived MSA queries and alignments underlying these depth bins, with templates disabled and protein-chain-only inputs. Confidence selects among five AF2 pTM models (three recycles), 25 AF3 samples (five seeds, ten recycles), or 25 Boltz-2 samples (one seed, ten recycles). These are controlled shared-MSA runs, not the models' full default search pipelines.




05_folding_iddt.png

The 248B-token model through Helico: request the top-L predicted contacts, generate three diffusion samples, select by highest Helico confidence at that fixed cut; no MSA. Helico drops pairs closer than six residues after coordinate mapping; requested and effective counts are saved. Same 305 natural proteins in every arm (95 validation + 210 test). GDT-TS and IDDT are separate metrics. See coverage.csv for exclusions. AlphaFold2/3 and Boltz-2 use the same archived MSA queries and alignments underlying these depth bins, with templates disabled and protein-chain-only inputs. Confidence selects among five AF2 pTM models (three recycles), 25 AF3 samples (five seeds, ten recycles), or 25 Boltz-2 samples (one seed, ten recycles). These are controlled shared-MSA runs, not the models' full default search pipelines.




06_sampling.png

Exp277 step 266,344, 314 natural proteins (97 validation + 217 test), 100 iid rollouts per protein. Consensus and oracle use the same sample pool within each protein. Individual maps are ranked by emission order; short maps retain denominator R, and malformed or unfinished individual maps score zero. The diagnostic consensus uses all parsed maps, following exp321. Oracle selection uses ground truth. This measures contact accuracy, not distinct folds. Scatter colors identify MSA tiers in the interactive view.



06_sampling_pl5.png

P@L/5 counterpart of the sampling diagnostic: the exact same 314 natural proteins and first 100 iid rollouts per protein. L is frozen input sequence length; $k = \max(1, \text{floor}(L/5))$. Individual maps retain emission order after eligibility filtering and deduplication; unfilled ranks receive zero credit and invalid/unfinished maps score zero. Consensus ranks votes from all parsed maps, matching the original diagnostic. Oracle best-of-100 is reselected using P@L/5 and uses ground truth. The scatter pairs consensus with this newly selected oracle per protein. Main contact-panel votes instead omit unfinished maps. This measures contact accuracy, not distinct folds.



06_sampling_pl5_long.png

P@L/5 counterpart of the sampling diagnostic: the exact same 314 natural proteins and first 100 iid rollouts per protein. L is frozen input sequence length; $k = \max(1, \text{floor}(L/5))$. Individual maps retain emission order after eligibility filtering and deduplication; unfilled ranks receive zero credit and invalid/unfinished maps score zero. Consensus ranks votes from all parsed maps, matching the original diagnostic. Oracle best-of-100 is reselected using P@L/5 and uses ground truth. The scatter pairs consensus with this newly selected oracle per protein. Main contact-panel votes instead omit unfinished maps. This measures contact accuracy, not distinct folds.

