

Question

Can FoldBench PPI support a useful experimental-complex benchmark under a 30% identity threshold?

We compare a strict component-held-out claim with a pair-held-out claim, then freeze a FoldBench-only set released after the inherited Protenix v1 training cutoff.

Two holdout claims

A qualifying alignment has exact identity $\geq 30\%$ over $\geq 50\%$ of the shorter sequence.

Component-held-out rejects a dimer when either chain has any training homolog. Pair-held-out rejects it only when one complex-training document has two chain instances that can be assigned one-to-one to homologs of both candidate chains.

The pair rule tests a new relationship between chain families. It permits prior single-chain exposure.

Survival

The strict rule leaves 0/184 scoped FoldBench dimers and 1/1,929 PINDER dimers after locally auditable MarInFold arms; the last PINDER target is present in Helico fine-tuning.

The pair rule leaves ****35 FoldBench + 183 PINDER = 218**** after both MarInFold complex arms. These include 55 homodimers and 163 heterodimers.

The conservative Helico same-PDB pair screen leaves ****32 + 65 = 97**** candidates.

Confirmation and evidence

The broad complex search hit its 100,000-result median cap. We repeated all 220 initial survivors at a 1,000,000-result cap; its median was 86,868 and it found two additional PINDER exclusions.

Exact MMseqs backtraces provide integer identity counts. Commands, hashes, compact witnesses and final status tables are committed. Complete compressed alignments and logs are published in the public MarInFold HF bucket.

Decision

Use FoldBench under the explicit ****pair-held-out**** claim. Its release-date design protects the Helico comparison from inherited Protenix v1 pretraining exposure. We still remove three targets with a paired homolog in Helico's later fine-tuning pool.

PINDER is not part of the frozen benchmark.

Frozen FoldBench evaluation

The 35 MarFold-clean pairs become **30 natural dimers** after the three Helico fine-tuning exclusions and manual removal of two de novo binder targets.

The set contains 13 homodimers and 17 heterodimers in 26 connected 30%-identity groups. One five-target ubiquitin-related group is kept intact. The deterministic split has **8 development and 22 test targets**; both cuts have the same mean length (448.25 vs 448.27 residues).

Every target ships with its exact FoldBench assembly, canonical two-chain input, label-to-author residue map, resolved cross-chain candidate universe, and contacts-v1 pyconfind ground truth.

Contact-context eligibility

A full-context diagnostic ran 100 rollouts for every target. Seven homodimers had at least one rollout reach the checkpoint's complete 8,192-token context: 147/3,000 rollouts in total. These remain valid structural targets but cannot support the strict exp82 requirement that every rollout finish.

The primary contact and contact-conditioned structural comparison therefore uses **23 context-complete targets** in 19 homology groups. Its metadata-only split has **6 development / 17 test targets**, 6 homodimers and 17 heterodimers, and mean lengths of 361.33 versus 361.29 residues.

Evaluation contract

Contact evaluation ranks only resolved cross-chain residue pairs. R is the number of degree ≥ 0.001 pyconfind interface contacts; no within-chain sequence-separation filter is applied across chains. Contact-budget choices use the 6-target development cut only; the 17-target test cut is read once for the final result.



The structural bundle uses FoldBench's native layout for Helico/DockQ and retains the exact two scored chains even when a deposited assembly has extra copies. Model sample selection must use model confidence; ground-truth best-of-N remains an oracle diagnostic.

The complete 12.3 MB bundle is public at ``hf://buckets/open-athena/MarinFold/data/evals/exp350_foldbench_pair_holdout/v1``.

The context audit and 23-target contact table are public at ``hf://buckets/open-athena/MarinFold/data/evals/exp350_foldbench_pair_holdout/contact_eval_v1``.


foldbench_eval_freeze.png

The 35 Marifold pair-clean FoldBench dimers yield 30 natural structural targets, of which 23 complete all 100 contact rollouts inside the model's 8,192-token context. The 6/17 contact split is grouped by chain homology and balances mean total length.



pair_survival.png

The strict component-held-out rule rejects a dimer when either chain has a training homolog. The pair-held-out rule rejects it only when one complex training document contains one-to-one homologs for both candidate chains. The final column additionally applies a conservative same-PDB pair screen against Helico's documented fine-tuning pool.



survival.png

Natural protein-only dimers after scope/quality filters and cumulative sequence exclusion against four locally available exp343 training arms. Every constituent chain must avoid a $\geq 30\%$ identity alignment covering $\geq 50\%$ of the shorter sequence. The remaining PINDER target is an upper bound because the ProteinMPNN redesign corpora were not searched.

