

Question

Does 1,000-short-rollout consensus improve over 100 full rollouts? Default exp277 checkpoint, step 266344. Eval-val only: 97 natural proteins.

Comparison

Caps: 5, 10, 20, $\text{floor}(L/5)$, $\text{floor}(L/2)$ emitted contact statements. Common 1,000 resampled trajectories yield every cap; 100-prefix controls. Fresh 100-full-rollout baseline; fixed temperature 1, top-p .95, top-k off. Exact exp89 metrics, paired protein bootstrap intervals, timings and raw tokens.

Status

Contact-cutoff tests and a real-protein GPU smoke pass. Production submitted on one GB200; all 97 proteins are queued for evaluation. No accuracy conclusion yet.