Question and design

Does model width improve contact accuracy at the same one-epoch exposure? Baseline: exp277, native plus MPNN corpus, 24 layers. Hidden width 2048 -> 4096; MLP width 8192 -> 16384. Keep 32 query / 8 KV heads; head width 64 -> 128. Keep data, optimizer schedule, seeds and global batch fixed.

Architecture and smoke verified

5,861,984,256 parameters, versus exp277's 1,471,374,336 (3.984x). 12 local configuration/tokenizer/data checks passed. 64-H100 smoke: ten updates, train loss 8.56948, eval loss 8.03575. 2.821 seconds/update; native and HF checkpoints saved with tokenizer. Compressed smoke schedule was spiky; this is an execution check.

Current status

Full 266,345-update epoch submitted for 64 H100s in US-EAST-02A. At 15:11 UTC, queued for eight additional free H100s. Global batch 128 retained through gradient accumulation. Production uses the original 10% warmup and 1e-3 peak LR. Compute-only estimate: 8.7 days, plus validation/checkpoint/preemption overhead. Contact-accuracy results are pending.